

AI Exposure at Three Levels: Reconciling Occupation, Task and Skill Measures of the Labour Market

Conway

Contents

Extended summary Introduction Section 1: Three descriptions of one labour market Section 2: Measures collected at three levels Section 3: Reconciling the measures Section 4: What the instrument will bear Section 5: Realised wages and employment Section 6: Within occupations Conclusion References

Extended summary

Exposure to generative artificial intelligence is measured at three different levels of the labour market. Felten, Raj and Seamans (2021) score the fifty-two abilities of the Occupational Information Network, Eloundou, Manning, Mishkin and Rock (2023) score its task statements, and Srinivasan, Chen and Zakerinia (2024) score occupations. Applied work uses each at whatever level it happens to need, and the aggregation or disaggregation that carries a score to that level has not been validated.

This paper puts the measures on a common footing, and what follows concerns the measures at least as much as the labour market they describe.

The propagation from tasks to skills can be checked against an independent measurement at the skill level, and it discriminates in the right direction. Because Felten’s index resolves into ten separate AI applications, the check is not a single number: carrying the Eloundou task ratings through the Network’s activity crosswalk reproduces the ability-level scores at a correlation of 0.667 for reading comprehension, 0.579 for translation and 0.546 for language modelling, against 0.289 for real-time video games and 0.172 for image recognition. The three language applications take the top three places. Eloundou’s own human and machine raters agree at 0.591, so the propagation performs at the reliability of its source on the applications that concern it and well below on those that do not.

The capability measures are less independent than their agreement implies. Srinivasan’s automation score is the Eloundou construction applied to a later release of the Network with a later version of the same model, on the same rubric and weights, so their correlation of 0.833 across occupations establishes reproducibility across vintages rather than convergence between approaches.

Augmentation is mostly a transformation of automation. Srinivasan and co-authors define it as one minus a Herfindahl index over the shares of exposed and unexposed tasks, which is a quadratic in the exposed share, and a quadratic in the automation score accounts for 77 percent of its variance. Net of that, skill groups explain 0.209 of what remains and eleven of twenty-five

groups load on it, with social skills, engineering and technology, arts and humanities, work output and psychomotor abilities on the augmenting side. Business and management and mental processes do not survive the correction, falling from statistics near three to 0.4 and 1.2, so the apparent concentration of automation on cognitive and administrative work is the index restating an exposure share.

Realised wages cannot settle whether any of this has occurred. Across a balanced panel of 693 occupations running from 2012 to 2025, no skill group moved by more than 1.51 times its own year-to-year volatility after 2022, and the observed period sits inside the range produced by relabelling ordinary years. Every deviation from the pre-2020 path begins in 2020 or 2021, before large language models were released, so the pandemic and the technology cannot be told apart in these data.

Comparing tasks within an occupation avoids that problem, since nothing is compared across time, and there a task's rated exposure does predict the automation share of its observed use, at a standardised value of 3.70 against two thousand within-occupation shuffles. The rating that predicts is the machine's. Substituting the human rating of the same tasks gives nothing, and the gap is too large for attenuation, so the relationship may be two language models agreeing with one another about what counts as model-like work.

Four properties of the Network constrain all of this. Its importance ratings compress towards uniformity when converted into shares. Its descriptors are collinear enough that the effective dimension of the space is about a third of its nominal dimension. Its two descriptions of what an occupation involves agree at a median rank correlation of 0.33. And its occupational content is revised too slowly to register a job changing while keeping its name. It is used here because it is free, covers the whole economy, and supplies the vocabulary in which every measure above is expressed.

Introduction

The empirical literature on generative artificial intelligence and work rests on a small number of exposure measures. Those measures have been compared with one another and their disagreements noted, usually as a question of whether capability ratings or observed usage is the more informative. Less attention has been paid to the fact that they were collected at different levels of description, so that comparing them at all requires a translation nobody has examined.

The translation matters because the levels differ in what they can support. Wages are observed once per occupation, and there are 829 of them in the 2025 Occupational Employment and Wage Statistics, carrying a wage bill of 10.70 trillion dollars. There are 18,796 task statements, so a set of task values has more unknowns than facts and cannot be contradicted by the data producing it. There are 161 descriptors, and after grouping twenty-five, so a set of skill values can fail. Whether a quantity in this literature is testable depends on how many task statements the Network's analysts wrote, which is a fact about a survey instrument. Freund and Mann (2026) cluster roughly nineteen thousand tasks into thirty-eight categories before estimating, and that step is what renders their estimation feasible.

Section 1 sets out the three layers. Section 2 places each measure at the level where it was collected, Section 3 reconciles them, and Section 4 reports what the underlying data will support. Sections 5 and 6 apply the apparatus, first to fourteen years of wages and employment and then to tasks within occupations.

Section 1: Three descriptions of one labour market

An industry produces output, that output is produced by activities, and those activities draw on human capability, so occupations, tasks and skills describe one process at three distances. The premise runs from Autor, Levy and Murnane (2003) and Acemoglu and Autor (2011) through the skill literature descending from Katz and Murphy (1992). Freund and Mann (2026) give it a general equilibrium form in which occupations bundle tasks, workers hold task-specific skills, and automating a task raises the weight on the tasks that remain; estimating that model on individual worker data, they project that the return to social and non-routine manual skills rises and the return to analytical skills falls.

The three layers do different work. The wage bill is a budget, and quantities constructed as shares of a total return that total whatever the shares are, so a reconciliation to the national wage bill tests nothing. Tasks are where the technology has been observed, since every measure of generative AI was built by judging a capability against a description of work. Skills are where a claim can fail, being fewer in number than the occupations whose wages are observed.

The skill grouping used throughout is the third level of the Network’s own content model, which gives twenty-five named groups across abilities, skills, knowledge and work activities. Figure 1 shows the three layers and the measures entering at each.

Section 2: Measures collected at three levels

Felten, Raj and Seamans (2021) built their measure by surveying respondents on the relatedness of ten AI applications to the fifty-two abilities of the Network, then combining the resulting matrix with occupational ability requirements, and they publish the intermediate ability-level scores. Those abilities are fifty-two of the 161 descriptors used here, so the measure enters the skill layer without propagation, and all fifty-two join after one renaming between releases. The ten applications come from 2016-era progress benchmarks and include image recognition, real-time video games and instrumental track recognition, so the headline index measures exposure to artificial intelligence broadly. The application-by-ability matrix allows any single one to be isolated, and for the present technology the language modelling and reading comprehension columns are the relevant ones.

Eloundou, Manning, Mishkin and Rock (2023) rate 19,265 task statements on a four-point rubric separating tasks a language model can substantially accelerate from those it cannot, with separate ratings by human annotators and by GPT-4, of which 18,179 join the task set used here. This is the measure Freund and Mann employ. Reaching skills requires the Network’s crosswalk from task statements to detailed work activities, which roll up to the forty-one generalized

activities and thence, through the published crosswalks from skills and abilities to work activities, to the wider descriptor set. Their headline index correlates 1.0000 with the machine labels and 0.591 with the human ones, so it is the machine rating; Section 6 turns on that distinction.

Srinivasan, Chen and Zakerinia (2024) score automation and augmentation for 911 occupations, of which 910 join the Network by name, giving 785 six-digit occupations and covering 681 of the 693 in the panel used here. Both indices are built from task descriptions of Network release 25.1 classified by GPT-4o on the Eloundou rubric. Job postings enter their paper to validate the indices and never to construct them, so neither score contains labour market outcomes.

The Anthropic Economic Index classifies conversations with a deployed model against Network tasks and reports the composition of interaction modes for each. Its automation share is the sum of directive and feedback-loop modes and its augmentation share the sum of iteration, validation and learning, with a filtered residual completing one hundred percent, so the two correlate minus 0.588 by construction. Coverage is 20.9 percent of the tasks considered here, and the covered tasks are more exposed by 0.236 on a zero to one scale than those omitted, with the most used tenth accounting for three quarters of recorded usage. It observes usage rather than capability.

Section 3: Reconciling the measures

Two of the measures reach ability-level exposure by routes sharing no data. Felten and co-authors obtain it from a survey on applications and abilities; propagating the Eloundou ratings through the activity crosswalk obtains it from a language model reading task text.

Propagated exposure correlates 0.667 with reading comprehension, 0.579 with translation and 0.546 with language modelling, and 0.289 with real-time video games and 0.172 with image recognition. The ordering follows the closeness of each application to what a language model does, and the video game column serves as a placebo. Since Eloundou’s human and GPT-4 raters agree on 65.6 percent of task labels with numeric scores correlating 0.591, the propagation performs at the reliability of its source on the applications that concern it.

Variation survives the aggregation. Averaging eighteen thousand task ratings into twenty-five groups gives a coefficient of variation of 0.292 across a range from 0.045 to 0.513, with mental processes highest, foundational literacy and numeracy close behind, and physical abilities lowest.

The Srinivasan automation score correlates 0.833 with Eloundou exposure aggregated to occupations, and the reason is instructive. Their score applies the Eloundou rubric to Network release 25.1 using GPT-4o, weighting the two exposed categories one and one half and weighting each task by its importance, so the correlation measures reproducibility across data vintages and model versions. Several capability measures in this literature share a rubric or a source of task descriptions, and agreement among them should be read accordingly.

Augmentation requires the same care. Srinivasan and co-authors define it as one minus a Herfindahl index over the shares of exposed and unexposed tasks in an occupation, on the reasoning that a job combining automatable work with work that resists automation gains most. That index is a quadratic in the exposed share, and the automation score is a related summary

of the same classification, so a quadratic in automation accounts for 77 percent of the variance in augmentation across occupations.

Inverting the two occupational scores onto the twenty-five skill groups needs no employment weighting, since exposure is not a share of the wage bill, so the effective sample stays at the number of occupations. Skill groups account for occupational exposure with cross-validated fit of 0.712 on automation and 0.415 on augmentation, and the named groups outperform ten principal components on both. On the raw scores, ten of twenty-five groups show an augmentation loading distinguishable from their automation loading, stable across five hundred bootstrap resamples. Removing the quadratic and inverting the residual leaves eleven groups distinguishable, with skill groups accounting for 0.209 of the residual variation, and the two profiles correlate 0.604. Social skills, engineering and technology, arts and humanities, work output and psychomotor abilities hold up; business and management falls from a statistic of minus 3.1 to plus 0.4 and mental processes from minus 2.9 to minus 1.2. Figure 2 shows both profiles.

Section 4: What the instrument will bear

Importance in the Network is recorded on a five-point scale of which respondents use a narrow part, and converting such a rating into a share yields something close to uniform. In the median occupation the most important task rates 1.41 times the least important. Where the descriptor mix of a task is set to its occupational average, all 124,614 occupation by descriptor pairs containing more than one task carry a single distinct value, and the skill layer becomes a fixed linear image of one number per occupation. A descriptor proximity matrix built from occupational similarity places 98.5 percent of its 798,342 pairs above 0.90, with a coefficient of variation of 0.022, and an occupation's profile across the forty-one generalized work activities sits a total variation distance of 0.10 from a uniform profile. The Network's two descriptions of what a job involves agree at a median rank correlation of 0.33.

Valuing 17,383 tasks by allocating each occupation's wage bill in proportion to task importance produces values correlating 0.997 on ranks with dividing that wage bill by the occupation's number of tasks, with eighty-seven of the largest hundred tasks identical under the two procedures. Task frequency, published over seven bands, varies about fifteen times more than importance with a median ratio of 14.78, and substituting it lowers the correlation to 0.883.

Fitting occupational wages with descriptor intensities and evaluating on held-out occupations gives 0.554 for a single ordinal measure of required preparation, 0.700 for twenty principal components, 0.648 for the twenty-five named groups and 0.544 for all 161 descriptors. Employment weighting, which a wage bill decomposition requires, gives an effective sample of 121 against 820 occupations in the sense of Kish (1965), the twenty largest occupations carrying a third of employment between them.

Skill quantities are employment multiplied by intensity and involve no estimation. Price levels are not recoverable: eleven of twenty-five groups take a negative price in some year, and constraining prices to be non-negative pins twenty of twenty-five at exactly zero. Price changes are recoverable, and skill content explains 52 percent of the variation in occupational wage changes between 2022

and 2025.

Inference needs more care here than it usually receives. Comparing any two periods in these data returns between five and nine of the twenty-five groups apparently changing course when nothing has happened. Conventional standard errors are computed against the precision of a window’s average rather than against the movement of the series between ordinary years, and composition series are smooth enough that a shift worth 0.17 percent of a level per year produces a statistic of 8.2. The procedure adopted here estimates the quantity for every year transition and builds the reference distribution by relabelling which transitions are treated. Four results in the course of this work reached conventional significance and did not survive it.

Section 5: Realised wages and employment

The panel covers 693 occupations from 2012 to 2025, with a wage bill rising from 5.15 to 8.95 trillion dollars, assembled from the Occupational Employment and Wage Statistics and mapped across the 2018 revision of the occupational classification using the Network’s crosswalk, which raises coverage of the 2018 wage bill from 0.813 to 0.936.

Across 2022 to 2025 no skill group moved by more than 1.51 times its own year-to-year volatility. Five of twenty-five exceed two of their own standard errors, against a median of four across 286 relabellings, and quantities give a probability of 0.455 weighted by headcount.

A directional test against the Freund and Mann projection, with group membership assigned from their vocabulary before any outcome was examined, is not supported on wages and runs in the opposite direction, the returns to both social and analytical content having risen against their pre-period paths with analytical rising further. The social against analytical contrast in employment moves as projected, at a probability of 0.021 against all relabellings and 0.273 against consecutive triples, on the one specification carrying no pre-existing trend.

Figure 3 shows deviations from the pre-2020 path by year. On wages the residual runs at minus 0.005 in 2020 and reaches minus 0.025 by 2025; on employment the series takes an excursion of minus 0.040 in 2020 and remains positive from 2021. Large language models were released in late 2022, after both breaks. The pandemic sits between the pre-period and the treatment period and moves these series by more than the effect sought, so a before-and-after comparison in United States occupational data does not separate the technology from the recovery.

Section 6: Within occupations

The pandemic moved which occupations grew. A technology of this kind should also move what happens inside an occupation, and that comparison runs across tasks within a job rather than across periods, so nothing about 2020 can enter it.

Selection has to be handled first. The Anthropic index records only tasks a deployed model has already been used for, so asking whether exposure predicts usage would select on the outcome. Asking instead whether exposure predicts the composition of that usage, conditional on the task

being used at all, puts the selection on volume and the outcome on mode. The outcome used is the automation share of automation plus augmentation, which removes the filtered residual that makes the raw percentages compositional.

On 1,617 tasks across 349 occupations, with occupation means removed, a task's exposure predicts the automation share of its observed use at a slope of 0.109 and a standardised value of 3.70 against two thousand within-occupation shuffles. It survives controlling for usage and holds in three of the four releases, failing only in January 2026.

The result is probably about the measures rather than about work. Substituting Eloundou's human rating of the same tasks gives a slope of minus 0.008 and a standardised value of minus 0.38, which is too far from the machine result to be attenuation. If the human rating were the same signal carrying more noise, the correlation of 0.356 between the two ratings in this sample would still leave a slope near 0.039 and a standardised value near 1.8. On the 747 tasks where the two ratings differ by more than 0.4, entering both together gives 0.204 to the machine and 0.029 to the human, so the outcome follows the machine almost entirely.

Both sides of the relationship are language model outputs, since the Anthropic index classifies conversations using a model and the predictor is a model's reading of task text. These data cannot distinguish the machine being the more accurate rater from the two sharing method variance. The second reading would matter for a literature increasingly built on model-generated task ratings, and distinguishing them would require a task-level outcome that no model has touched.

Conclusion

Three literatures measure exposure to generative artificial intelligence at three levels of one labour market. This paper states them together on public data, checks the propagation from tasks to skills against an independent measurement and finds it discriminating correctly among AI applications, and finds the capability measures closer to a single construct than their agreement suggests, with augmentation largely a transformation of automation and a residual that still carries skill structure. It documents four respects in which the Network constrains what may be concluded. On realised wages the before-and-after design does not identify an effect, the pandemic break preceding the technological one in every series examined. Within occupations a relationship does appear, and it holds only for the machine-rated exposure, which is a result about the instruments as much as about the labour market.

References

- Acemoglu, D. and Autor, D. (2011). Skills, tasks and technologies: implications for employment and earnings. *Handbook of Labor Economics*, volume 4B.
- Autor, D., Levy, F. and Murnane, R. (2003). The skill content of recent technological change: an empirical exploration. *Quarterly Journal of Economics*, 118(4).

Eloundou, T., Manning, S., Mishkin, P. and Rock, D. (2023). GPTs are GPTs: an early look at the labor market impact potential of large language models.

Felten, E., Raj, M. and Seamans, R. (2021). Occupational, industry, and geographic exposure to artificial intelligence: a novel dataset and its potential uses. *Strategic Management Journal*, 42(12).

Freund, L. B. and Mann, L. F. (2026). Job transformation, specialization, and the labor market effects of AI. Working paper.

Katz, L. and Murphy, K. (1992). Changes in relative wages, 1963-1987: supply and demand factors. *Quarterly Journal of Economics*, 107(1).

Kish, L. (1965). *Survey Sampling*. Wiley.

Srinivasan, S., Chen, W. X. and Zakerinia, S. (2024). Displacement or complementarity? The labor market impact of generative AI. Harvard Business School Working Paper 25-039.